
Radio Networks

The Model

Broadcast


Andrea CLEMENTI

A **radio network** is a set of **stations** (nodes) located over a support Euclidean Space.


To each node v , a transmission range $R(v) > 0$ is assigned.

A node w can receive a msg M from v **only if**

$$d(v, w) \leq R(v)$$



When a node v sends a msg M , M is sent over all the disk (Broadcast Transmission) in one **TIME SLOT**




Radio Networks are SYNCHRONOUS SYSTEMS

All nodes share the same global clock. So,


Nodes act in **TIME SLOTS**

Message transmissions are completed

within **one time slot**



Andrea CLEMENTI




Andrea CLEMENTI

The **Range Assignment** uniquely determines a

Directed Communication Graph $G(V,E)$

All in-neighbors of **S** receive
the **msg** in 1 HOP


unless.....



MESSAGE COLLISIONS

If, during a time slot,
two or more in-neighbors send a msg to **v** THEN

v does not receive anything.



RADIO MODEL:

a node v receives a msg during time slot T


IFF

there is exactly one of its in-neighbors that sends a msg during time slot T

TASK:

BROADCAST OVER A RADIO NETWORK $G(V,E)$

NOTE: FLOODING DOES NOT WORK !!!!!



Andrea CLEMENTI

CORRECTNESS (Strongly-Conn. $G(V,E)$, source s) :

A Protocol **completes** Broadcast from s over G
if there is one time slot s.t. every node is **INFORMED**
about the source msg.

TERMINATION

A Protocol **terminates** if **there is** a time slot t s.t.
every node **stops** any action **WITHIN** time slot t .

HOW can we **AVOID** MSG COLLISIONS ???

IDEA: ROUND ROBIN !!!

Start with Assumptions:

- nodes know a *good apx* of $|V| = n$
- nodes are indexed by $0, 2, \dots, n-1$

then

IDEA 1: ROUND ROBIN !!!

Start with Assumptions:

- nodes know a *good apx* of $|V| = n$
- nodes are indexed by $0, 2, \dots, n-1$

then

ROUND ROBIN PHASE

A *Phase* of ROUND ROBIN consists of n time-slots

At TIME $T = 0, 1, 2, \dots$

- NODE $i=T$, if informed, sends the source msg;
- All the Others do NOTHING

What can we say **AFTER** one Phase of **RR** ?

Assume that $\text{label}(s) = J$ (initially
J is the only informed one)

During the *FIRST PHASE* (*n time slots*):

Fact: ALL out-neighbors of s will be informed
after the First PHASE.

No MSG Collision occurs...

IDEA 2:


LET'S RUN THE RR PHASE FOR L consecutive times

THM. After Phase k , All nodes within Hop-Distance k from the source s

Proof. By induction on $HOP-DISTANCE = PHASE$ k

Inductive Step: Phase k

Informed Nodes



at time slot j :

- j sends to all its out-neighbors w
- no others are active

So, ALL w 's will receive the msg.

This Argument holds for all nodes in $L(k-1)$.

So all nodes in $L(k)$ will be informed after Phase k

Corollary (**RR COMPLETION TIME**).

Let D be the (unknown) **source** eccentricity. Then,

D **RR-Phases** suffice to INFORM all NODES

WHAT ABOUT TERMINATION ???

... It depends on the Knowledge of Nodes.

If they know **n** they CAN decide to stop... !

WHEN ????

The (unknown) **source eccentricity** is at most **$n-1$** ,

so....

They all have the global clock $\Rightarrow$ they all can decide to **stop** AFTER the **RR** Phase **$n-1$**

THM. Protocol **RR**

- **completes** Broadcast in **$D \times n$**
- **terminates** Broadcast in **$O(n^2)$**

Terrible question.....

What can we say if

NODES DO NOT KNOW any good bound on **n**

????

COMPLEX RESULTS:

-In UNKNOWN RADIO NETWORKS,

RR **Completes** in $O(D \cdot n) = O(n^2)$ time slots
Termination ?????

-There is an **optimal** Protocol that **completes**
in

$O(n \log^2 n)$ time slots

OBS.

RR does not exploit **parallelism** *at all*

GOAL:

SELECT PARALLEL TRANSMISSIONS

A “selective” method.

DEF. Given $[n] = \{1, 2, \dots, n\}$ and $k \leq n$,
a family of subsets

$$H = \{H_1, H_2, \dots, H_t\}$$

is (n, k) -selective if for any subset $S \subset [n]$ s.t.
 $|S| \leq k$, an $H \in H$ exists s.t.

$$|S \cap H| = 1$$

Trivial Fact.

The family $H = \{\{1\}, \{2\}, \dots, \{n\}\}$ is (n, k) -selective for any k .

How a selective family can be used to BROADCAST?

Restriction: Nodes know n and d ;

(** As for the *completion* time: they can be removed)

SET UP:

All nodes know the same (n,d) -selective family

$$H = \{H_1, H_2, \dots, H_i, \dots, H_t\}$$

where

$$d = \text{max-degree}(G)$$

Protocol SELECT1.

- Protocol works in consecutive Phases $J=1,2,\dots$ (as **RR** !!!).
- At time slot i of every Phase,
every informed node in H_i transmits

Protocol Analysis.

- **Lemma 1.** After Phase j , all nodes at distance at most j will be informed.

Proof. By induction on j . $j=1$ is trivial. Then, consider a node y at distance j . Consider the node subset

$$N(y) = \{z \in V \mid z \text{ is a neighbor of } y \text{ \& } z \text{ is at distance } j-1\}$$

Since $N(y) \subseteq [n]$ and $|N(y)| \leq d$, apply (n,d) -selectivity and get the thesis.

Is it correct?

NO!!!!

We are not considering the impact of informed nodes z in level j during phase j !

a) if you put z into $N(y)$, z could be selected but not already informed

b) if you don't put z into $N(y)$, z could be informed and create collisions

So what?

A very simple **change** makes the protocol correct!!!

ONLY NODES THAT HAVE BEEN INFORMED
DURING PHASE $j-1$
WILL BE **ACTIVE** DURING PHASE J

No unpredictable collisions and enough to inform
level j

Lemma 1 is now true!, so after D phases, all levels will be informed.

Completion time is $O(D |H|)$

So we need **minimal-size selective families**.

THM (ClementiMontiSilvestri 01).

For sufficiently large n and $k \leq n$, there exists an **(n,k) -selective** family of size

$O(k \log n)$ and this is optimal !

If we plug-in the minimal size (n,d) -selective family into the protocol, we get:

$O(D d \log n)$ time

So if D and d are both *small* (most of "good" networks), we have a *much better time* than the **RR** one

THE LOWER BOUND.

Can the selective protocol be improved for general graphs?

NO!

THM. In **directed general graphs**, the use of a **selective family** is somewhat **necessary**,

GET for $Dd \leq n$: $\Omega(D d \log(n/D))$

LOWER BOUND.

Construct a Layered Directed Network.

$L_0 = \{s\}$, then L_j as follows:

Let $m < \min \text{size } (n/D, d)\text{-selective family}$.

Adv chooses the next level by looking at Prot's transmissions for the next m time slots as if L_j was ALL the rest of nodes.

He then chooses the subset of nodes not selected by Prot (since $m < \min \text{size } (n/D, d)\text{-selective}$).

This subset becomes L_j

OBS.


- *Adv* can do this for $O(n/D)$ levels in order to produce a network of diameter D still keeping $|R| > n/2$.

- The behaviour of *Prot* is the same in both scenarios:

R = ALL THE REST OF NODES

R = LJ

THE LOWER BOUND (Proof).



Bipartite Complete Graph between L_{j-1} and the unselected subset of R

Andrea CLEMENTI

Proof (LOWER BOUND).

-The Layered Graph shows that, in order to inform each Level, Prot needs to produce a transmission scheduling

$$H = \{H_1, \dots, H_k\}$$

which must be

$(n/D, d)$ -selective.

So $|H|$ must be $\Omega(d \log(n/D))$ and globally get

$\Omega(D * d \log(n/D))$ time.

Random vs Deterministic: an **Exponential** Gap

Lower Bound for deterministic protocol when

$$d = n \text{ and } D = 3 \quad \rightarrow \quad \Omega(n \log n)$$

What about **Randomized** Protocols ?

Example: at every time slot,

every **informed** node transmits with probability **1/2**.

The **BGI** RND Protocol
(Case of **d**-regular layered graphs (as in the L.B))

Repeat for **K** = 1,2,... (Stage)

Repeat for **j** = 1,2, ..., **c log n**

If node **x** has been *informed* in Stage **k-1**
then **x** transmits with probability

$$1/d$$

Protocol Analysis.


THM. Prot. BGI completes Broadcast within
 $O(D)$ Stages, so within

$O(D \log n)$ time step

WITH HIGH PROBABILITY

PROOF. By Induction on Level $L=1\dots D$.

$D=1 \rightarrow$ Trivial. So assume all nodes of L_j are informed after $t = O(j \log n)$ time slots. Consider STAGE $j+1$.



Which is the Prob
that y will be informed
during STAGE $J+1$?

Probability in 1 time slot:

$$d * (1/d) (1-1/d)^{d-1} > 1/8$$

Probability that he is **not** informed
in (1 Stage =) $c \log n$ independent time slots:

$$< (1-1/8)^{c \log n} < e^{-c/8 \log n} < 1/n^{c/8}$$

since

- **Independent *rnd* choices**
- $(1-x) < e^{-x}$ for any $0 < x < 1$

we need this for *all* nodes ($< n$)

apply **UNION BOUND** *twice*:

$$* \Pr(\exists \text{ BAD node }) < n (1/n^{\{c/8\}}) < 1/n^{\{c-1/8\}}$$

we need this for $k = D < n$ Stages

$$** \Pr(\exists \text{ BAD Stage }) < 1/n^{\{c-2/8\}}$$

By choosing $c > 10$, you get Theorem

WITH HIGH PROBABILITY = $(1-1/n)$

(*) Task:

Extend the BGI Protocol to

General Graphs

So to complete Broadcast in

$O(D \log^2 n)$ time slot (W.H.P.)

Restriction: nodes know n

You are interesting in [learning](#) more?

See the paper (CMS01.pdf)
in the Course Web Page

Thanks!

Andrea